Exercise 7

Deadline: 23.1.2019

Regulations

Please hand in the following:

- Jupyter notebook and exported HTML `regularized-regression-commented.*` for question 1.
- PDF (hand written or typed) `proof-ridge-regression.pdf` for question 2.
- Jupyter notebook and exported HTML `kernelized-ridge-regression.*` for question 3.
- Jupyter notebook and exported HTML `fitting-circles.*` for question 4.

Zip all files into a single archive with naming convention (sorted alphabetically by last names)

`lastname1-firstname1_lastname2-firstname2_exercise06.zip`

or (if you work in a team of three)

`lastname1-firstname1_lastname2-firstname2_lastname3-firstname3_exercise06.zip`

and upload it to Moodle before the given deadline.

1 Comment on your solution to exercise 6

Study the sample solution `ex06-solution.ipynb` to be provided on Moodle and use it to comment on your own solution to this exercise. Specifically, copy your original notebook `regularized-regression.ipynb` to `regularized-regression-commented.ipynb` and export it to `regularized-regression-commented.html` in the end. Insert comments as markdown cells starting with

`<span style="color:green;font-weight:bold">Comment</span>`

in order to clearly distinguish your comments from other cell types. The point of these comments is that you identify your errors and bugs yourselves, so that you learn from your mistakes. In addition, the tutor will have an easier time distinguishing between the first mistake and consequential errors caused by the first one and will only deduct points for the former. If you fail to hand in comments, the tutor is not required to make this distinction and will deduct points for all errors alike.

2 Proof - Ridge Regression - Primal vs Dual (10 Points)

In the primal formulation, the ridge regression problem takes the following form:

$$\widehat{\beta} = \text{argmin}_\beta \left\| \mathbf{y} - X\beta \right\|_2^2 + \tau \left\| \beta \right\|_2^2 , \tag{1}$$

where X is an $N \times D$ matrix, β is a D-dimensional vector and $\mathbf{y}$ is an N-dimensional vector. As you saw in the lecture, the optimal $\widehat{\beta}$ is given by

$$\widehat{\beta} = \left(X^T X + \tau \mathbb{1}_D \right)^{-1} X^T \mathbf{y} . \tag{2}$$

Here $\mathbb{1}_D$ is the D-dimensional unit matrix. You also know that the dual formulation of the problem is given by

$$\widehat{\alpha} = \text{argmax}_\alpha - \alpha^T \left(X X^T + \tau \mathbb{1}_N \right) \alpha + 2\alpha^T \mathbf{y}, \tag{3}$$

with the solution for $\widehat{\alpha}$

$$\widehat{\alpha} = \left(X X^T + \tau \mathbb{1}_N \right)^{-1} \mathbf{y} . \tag{4}$$

For each feasible α, a corresponding feasible β is given by:

$$\beta = X^T \alpha \,. \tag{5}$$

Prove that the optimal $\widehat{\beta}$ corresponds to the optimal $\widehat{\alpha}$:

$$\widehat{\beta} = X^T \widehat{\alpha} \,. \tag{6}$$

Hint: Prove the following lemma, which will be useful in your derivation:

$$\left(X^T X + \tau \mathbb{1}_D\right)^{-1} X^T = X^T \left(X X^T + \tau \mathbb{1}_N\right)^{-1} \tag{7}$$

3 Kernelized (Ridge) Regression (10 Points)

From the lecture you know that using the dual formulation of the ridge regression, we are able to introduce the **kernel trick** with kernel matrix G. New regressed values y_{new} are computed through

$$\mathbf{y}_{new} = \mathbf{y}^T \left(G + \tau \mathbb{1}_N\right)^{-1} \kappa^T \,. \tag{8}$$

The elements of the kernel matrix are computed on the training set via the kernel function K

$$G_{i_1 i_2} = K(X_{i_1}, X_{i_2}), \tag{9}$$

and the weight vector $\alpha = \mathbf{y}^T \left(G + \tau \mathbb{1}_N\right)^{-1}$ only needs to be computed once at the beginning. In contrast, the kernel vector κ has to be recomputed for each new instance X_{new} according to $\kappa_i = K(X_{\text{new}}, X_i)$.

In this task, you shall apply kernel ridge regression to reconstruct all missing pixels in the grayscale image `cc_90.png` (to be found on Moodle). The features X_i are the pixel coordinates, and the response y_i is the corresponding grayvalue. Pixels with grayvalue = 0 are considered missing and shall be replaced with their regressed values. Use a squared exponential kernel function

$$K(X_{i_1}, X_{i_2}) = \exp\left(-\frac{\|X_{i_1} - X_{i_2}\|^2}{2\sigma^2}\right) \tag{10}$$

If you are careless in structuring your program, it will be too slow to run. Therefore, keep to the following guidelines:

- Cut off the kernel function (i.e. set it to zero) at a sensibly large radius to get a sparse kernel matrix, and use a sparse array class.
- Do not loop over each cell in the kernel matrix, vectorize at least one dimension.
- Do not explicitly invert any matrices. In general, but especially for sparse matrices, solving $A\mathbf{x} = b$ for $\mathbf{x}$ is many times faster than computing $\mathbf{x} = A^{-1}\mathbf{b}$ explicitly. For example, you can use `scipy.sparse.linalg.spsolve`

Do the same experiment for the Nadaraya-Watson kernel regression (no ridge):

$$\mathbf{y}_{new} = \frac{\sum_i \mathbf{y}_i \kappa_i}{\sum_i' \kappa_i'} \tag{11}$$

Play with the σ of the squared exponential as well as the τ of the ridge regression and find the parameters that produce the visually best reconstructed image for both approaches. Comment on the quality of the two methods.

4   Fitting Circles

In this exercise we will have a closer look at fitting circles to data. The numpy-file `circles.npy` found on Moodle contains many pairs of x-y-coordinates, and can be loaded through `data = np.load("circles.npy")`. Visualize the data in a scatter plot to show that the points are arranged in the shape of several circles and circle segments. Pay attention that the axes are scaled identically when plotting the data, otherwise your circles will look like ellipses. How many circles or circle segments would you fit into the data as a human?

4.1   RANSAC (6 Points)

First implement the RANSAC-Algorithm for the fitting of circles:

- For a set number of times N, repeat the following:
  - Randomly choose 3 points and determine the circle passing throuhg all of them, parametrized though its radius and the coordinates of the center. For example, you can achive this by foming and solving a system of three equations, one for each point.
  - Classify points that are closer than ϵ to the circle as inliers and count them.
  - If the inlier count for this circle is higher than for the best circle so far, save the current circle and its inliers as the new best.

- If you want to fit further circles, delete all inliers of the last fitted circle from the dataset and repeat the procedure.

Then, estimate the number of iterations N that is needed, as you learnt in the lecture. Plot all fitted circles on top of original data and comment on the result. Is the result sensitive to the value of ϵ, the parameter that is used to determine the inliers?

Hint: For plotting circles, you can use the following methods:

```
circle = plt.Circle((cx, cy), radius=r, fill=False) # Create a circle
plt.gca().add_patch(circle) # Add it to the plot
```

In the next two sub-tasks, you will further improve the fits of your circles using two different methods. If you were not able to implement the RANSAC, then either try to fit a circle manually in the data and get the inliers this way, or create data for a circle + noise and use it for the rest of the exercise.

4.2   Algebraic Distance (4 Points)

For each set of inliers you found with RANSAC, fit a circle by minimizing the algebraic distance:

$$\min_{c,r} \sum_i \left(||x_i - c||_2^2 - r^2 \right)^2 \tag{12}$$

To do this, reformulate (12) as $\widehat{\beta} = \operatorname{argmin}_\beta \sum_i (\tilde{Y} - \tilde{X}_i \beta)^2$, following the transformations from the lecture, and solve using the usual least-squares method. Plot the refined circles on top of the data and comment on the results.

4.3   Levenberg-Marquardt (6 Points)

For each set of inliers you found with RANSAC, solve

$$\min_{c,r} \sum_i \left(||x_i - c||_2 - r \right)^2 \tag{13}$$

by adapting the Levenberg-Marquardt algorithm from the lecture to this problem. Derive the theoretical approach by hand, but feel free to use `scipy.optimize` (or another optimization library of your choice) as an implementation.

For the theory, first solve for r and express it in terms of x, c, then calculate the derivative with respect to c. Again, plot the resulting circles and comment on what you find.

4.4 Comparison (4 Points)

It was mentioned in the lecture that the solution from (13) is more robust to outliers than solving (12). Try to show this experimentally.

Hint: You can approach this problem for example as follows: Pick one set of inliers, sample a subset of those and maybe add a couple of outliers. Then apply your procedures from 3.2 and 3.3 to this dataset.